

## Q6 Batch Norm and SELU (4.5 Points)

Generate training and test datasets for a binary classification problem using Fashion-MNIST with class 1 being a combination of sneaker and pullover and class 0 being the combination of sandal and shirt categories.

- Train the model using Logistic regression (No Hidden Layers). Report train and test loss.
- Train a Neural Network with one hidden layer (100 neurons). Use Adam optimizer and Relu activation for hidden layer. First overfit a small sample to check errors and get idea of learning rate. Then train on complete dataset. Add regularization (dropout or weight decay) if needed.
- Now add another hidden layer (50 Neurons). Adjust the learning rate if you have to. Add regularization (dropout or weight decay) if needed.
- Now try adding Batch Normalization and compare the train and test loss : Is it converging faster than before? Does it produce a better model? How does it affect training speed? **Do not use dropout with batch normalization.**
- Try replacing Batch Normalization with SELU, and make the necessary adjustments to ensure the network self-normalizes (i.e., standardize the input features, use LeCun normal initialization, make sure the DNN contains only a sequence of dense layers). Compare the results with Batch Normalization. **For SELU if you are using dropout then use alpha dropout.** Alpha dropout make sure that network is self normalized.

### ▼ Importing Libraries

```
if 'google.colab' in str(get_ipython()):
    print('Running on Colab')
else:
    print('Not running on Colab')
```

Running on Colab

```
if 'google.colab' in str(get_ipython()):
    !pip install wandb --upgrade
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-whee
Collecting wandb
  Downloading wandb-0.13.4-py2.py3-none-any.whl (1.9 MB)
    |████████████████████████████████████████| 1.9 MB 19.2 MB/s
Collecting sentry-sdk>=1.0.0
  Downloading sentry_sdk-1.10.1-py2.py3-none-any.whl (166 kB)
    |████████████████████████████████████████| 166 kB 74.2 MB/s
Requirement already satisfied: protobuf!=4.0.*,!=4.21.0,<5,>=3.12.0 in /usr/local
Requirement already satisfied: six>=1.13.0 in /usr/local/lib/python3.7/dist-packa
Collecting pathtools
```

```

Downloading pathtools-0.1.2.tar.gz (11 kB)
Requirement already satisfied: Click!=8.0.0,>=7.0 in /usr/local/lib/python3.7/dist-packages (7.0)
Collecting GitPython>=1.0.0
  Downloading GitPython-3.1.29-py3-none-any.whl (182 kB)
    |████████████████████████████████████████| 182 kB 60.8 MB/s
Collecting shortuuid>=0.5.0
  Downloading shortuuid-1.0.9-py3-none-any.whl (9.4 kB)
Requirement already satisfied: psutil>=5.0.0 in /usr/local/lib/python3.7/dist-packages (5.9.0)
Requirement already satisfied: PyYAML in /usr/local/lib/python3.7/dist-packages (6.0)
Collecting setproctitle
  Downloading setproctitle-1.3.2-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux2014_x86_64.whl (25 kB)
Requirement already satisfied: setuptools in /usr/local/lib/python3.7/dist-packages (57.5.0)
Requirement already satisfied: requests<3,>=2.0.0 in /usr/local/lib/python3.7/dist-packages (2.28.1)
Collecting docker-pycreds>=0.4.0
  Downloading docker_pycreds-0.4.0-py2.py3-none-any.whl (9.0 kB)
Requirement already satisfied: promise<3,>=2.0 in /usr/local/lib/python3.7/dist-packages (2.3.0)
Collecting gitdb<5,>=4.0.1
  Downloading gitdb-4.0.9-py3-none-any.whl (63 kB)
    |████████████████████████████████████████| 63 kB 2.3 MB/s
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.7/dist-packages (4.4.0)
Collecting smmap<6,>=3.0.1
  Downloading smmap-5.0.0-py3-none-any.whl (24 kB)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (3.4)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (2022.9.24)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (3.7.4)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (1.26.13)
Collecting sentry-sdk>=1.0.0
  Downloading sentry_sdk-1.10.0-py2.py3-none-any.whl (166 kB)
    |████████████████████████████████████████| 166 kB 65.7 MB/s
  Downloading sentry_sdk-1.9.10-py2.py3-none-any.whl (162 kB)
    |████████████████████████████████████████| 162 kB 28.5 MB/s
  Downloading sentry_sdk-1.9.9-py2.py3-none-any.whl (162 kB)
    |████████████████████████████████████████| 162 kB 81.1 MB/s
  Downloading sentry_sdk-1.9.8-py2.py3-none-any.whl (158 kB)
    |████████████████████████████████████████| 158 kB 79.9 MB/s
  Downloading sentry_sdk-1.9.7-py2.py3-none-any.whl (157 kB)
    |████████████████████████████████████████| 157 kB 77.2 MB/s
  Downloading sentry_sdk-1.9.6-py2.py3-none-any.whl (157 kB)
    |████████████████████████████████████████| 157 kB 70.7 MB/s
  Downloading sentry_sdk-1.9.5-py2.py3-none-any.whl (157 kB)
    |████████████████████████████████████████| 157 kB 63.5 MB/s
  Downloading sentry_sdk-1.9.4-py2.py3-none-any.whl (157 kB)
    |████████████████████████████████████████| 157 kB 73.7 MB/s
  Downloading sentry_sdk-1.9.3-py2.py3-none-any.whl (157 kB)
    |████████████████████████████████████████| 157 kB 64.2 MB/s
  Downloading sentry_sdk-1.9.2-py2.py3-none-any.whl (157 kB)
    |████████████████████████████████████████| 157 kB 64.8 MB/s

```

```

# mount google drive
if 'google.colab' in str(get_ipython()):
    from google.colab import drive
    drive.mount('/content/drive')

```

Mounted at /content/drive

```

# Importing the necessary libraries
import torch

```

```

import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F

from torch.optim.lr_scheduler import ReduceLROnPlateau, ExponentialLR, CyclicLR, OneCycleLR

import numpy as np
import random

from datetime import datetime
from pathlib import Path
import sys
from types import SimpleNamespace

import wandb

```

## ▼ Initializing Wandb

```
wandb.login()
```

```

ERROR:wandb.jupyter:Failed to detect the name of this notebook, you can set it manually
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
True

```

```
wandb.init(name = "Hw5_FASHION_MNIST.ipynb", project = 'Deep_Learning_Class_UTD')
```

```

wandb: Currently logged in as: pranavshekhar2. Use `wandb login --relogin` to force relogin
Tracking run with wandb version 0.13.4
Run data is saved locally in /content/wandb/run-20221028_214025-23f1zng3
Syncing run Hw5_FASHION_MNIST.ipynb to Weights & Biases (docs)

```

## ▼ Q6 Batch Norm and SELU (4.5 Points)

Generate training and test datasets for a binary classification problem using Fashion-MNIST with class 1 being a combination of sneaker and pullover and class 0 being the combination of sandal and shirt categories.

- Train the model using Logistic regression (No Hidden Layers). Report train and test loss.
- Train a Neural Network with one hidden layer (100 neurons). Use Adam optimizer and Relu activation for hidden layer. First overfit a small sample to check errors and get idea of learning rate. Then train on complete dataset. Add regularization (dropout or weight decay) if needed.

- Now add another hidden layer (50 Neurons). Adjust the learning rate if you have to. Add regularization (dropout or weight decay) if needed.
- Now try adding Batch Normalization and compare the train and test loss : Is it converging faster than before? Does it produce a better model? How does it affect training speed? **Do not use dropout with batch normalization.**
- Try replacing Batch Normalization with SELU, and make the necessary adjustments to ensure the network self-normalizes (i.e., standardize the input features, use LeCun normal initialization, make sure the DNN contains only a sequence of dense layers). Compare the results with Batch Normalization. **For SELU if you are using dropout then use alpha dropout** Alpha dropout make sure that network is self normalized

```
# This is the path where we will download and save data
if 'google.colab' in str(get_ipython()):
    data_folder = Path('/content/drive/MyDrive/Deep_Learning_UTD/Dataset')
    model_folder = Path('/content/drive/MyDrive/Deep_Learning_UTD/Model')
else:
    data_folder = Path('/home/harpreet/Insync/google_drive_shaanmoor/data/datasets')
```

## ▼ Train and Test Dataset - FASHION\_MNIST

```
# Transform to convert images to pytorch tensors
trans1 = transforms.ToTensor()

# Transform to normalize the data
# The mean and std are based on train subset which we will create below
trans2 = transforms.Normalize((0.2857,), (0.3528))
trans = transforms.Compose([trans1, trans2])

# Download the training_validation data (we will create two subsets - trainset and valset
train_val_set = torchvision.datasets.FashionMNIST(root = data_folder,
                                                  train = True,
                                                  transform = trans,
                                                  download = True)

# Download the testing data
testset = torchvision.datasets.FashionMNIST(root = data_folder,
                                             train = False,
                                             transform = trans,
                                             download = True)

def split_dataset(base_dataset, fraction, seed):
    split_a_size = int(fraction * len(base_dataset))
    split_b_size = len(base_dataset) - split_a_size
    return torch.utils.data.random_split(base_dataset, [split_a_size, split_b_size], gener
)

trainset, validset = split_dataset(train_val_set, 0.8, 10)
```

```
#DataLoaders

# Initializing the batch size
batch_size = 256

# Creating data loader for train set
train_loader = torch.utils.data.DataLoader(dataset= trainset,
                                             batch_size = batch_size,
                                             shuffle = True)

valid_loader = torch.utils.data.DataLoader(dataset = validset,
                                             batch_size = batch_size,
                                             shuffle = False)

# Creating data loader for test set
test_loader = torch.utils.data.DataLoader(dataset = testset,
                                             batch_size = batch_size,
                                             shuffle = False)

#Checking the mapping of the labels

train_val_set.data
train_val_set.class_to_idx

{'T-shirt/top': 0,
 'Trouser': 1,
 'Pullover': 2,
 'Dress': 3,
 'Coat': 4,
 'Sandal': 5,
 'Shirt': 6,
 'Sneaker': 7,
 'Bag': 8,
 'Ankle boot': 9}

#Creating a subset of data consisting of class 1 and class 0

from torch.utils.data.dataset import Subset

#Sneaker and Pullover - Class1
train_idx_1 = np.where((train_val_set.targets==7) | (train_val_set.targets==2))[0]
train_subset_1 = Subset(train_val_set,train_idx_1)

#Sandal and Shirt - Class 0
train_idx_2 = np.where((train_val_set.targets==5) | (train_val_set.targets==6))[0]
train_subset_2 = Subset(train_val_set,train_idx_2)

#Sneaker and Pullover - Class 1 - Creating labels for Class 1 Group

train_subset_1.dataset.data[train_subset_1.indices].shape
```

```

labels_subset_1 = [1]*(train_subset_1.dataset.data[train_subset_1.indices].shape[0]) #Label

train_val_set.targets[train_subset_1.indices]

        tensor([2, 7, 2, ..., 2, 7, 2])

#Sandal and Shirt - Class 0 - Creating labels for Class 0 Group

train_subset_2.dataset.data[train_subset_2.indices].shape
labels_subset_2 = [0]*(train_subset_2.dataset.data[train_subset_2.indices].shape[0])

train_val_set.targets[train_subset_2.indices]

        tensor([5, 5, 5, ..., 6, 5, 5])

train_data1 = []
for i in range(len(train_subset_1)):
    train_data1.append([train_subset_1[i], labels_subset_1[i]])

train_data2 = []
for j in range(len(train_subset_2)):
    train_data2.append([train_subset_2[j], labels_subset_2[j]])

#Combined Dataset
train_data1.extend(train_data2)

```

## ▼ Class for Model

```

class FashionMNIST_NN(nn.Module):
    def __init__(self, input_dim, output_dim, h_sizes, dprob, non_linearity, batch_norm):
        super().__init__()
        self.input_dim = input_dim
        self.h_sizes = h_sizes # list of hidden sizes
        self.non_linearity = non_linearity
        self.batch_norm = batch_norm
        self.dprob = dprob # list of dropout probabilities
        self.output_dim = output_dim

        # Initialize hidden layers
        model_layers = [nn.Flatten()]

        # hidden layers
        for i, hidden_size in enumerate(self.h_sizes):
            model_layers.append(nn.Linear(input_dim, hidden_size))
            model_layers.append(self.non_linearity)
            model_layers.append(nn.Dropout(p=dprob[i]))

        if self.batch_norm:
            model_layers.append(nn.BatchNorm1d(hidden_size, momentum=0.9))

```

```

        input_dim = hidden_size

    # output layer
    model_layers.append(nn.Linear(self.h_sizes[-1], self.output_dim))

    self.module_list = nn.ModuleList(model_layers)

def forward(self, x):
    for layer in self.module_list:
        x = layer(x)

    # we are not using softmax function in the forward pass
    # nn.crossentropy loss (which we will use to define our loss) combines nn.LogSoftmax(
    return x

def train(train_loader, loss_function, model, optimizer, grad_clipping, max_norm, log_batch_
    # Training Loop

    # initialize variables as global
    # these counts will be updated every epoch
    global batch_ct_train

    # Initialize train_loss at the he start of the epoch
    running_train_loss = 0
    running_train_correct = 0

    # put the model in training mode

    model.train()
    # Iterate on batches from the dataset using train_loader
    for input_, targets in train_loader:

        # move inputs and outputs to GPUs
        input_ = input_.to(device)
        targets = targets.to(device)

        # Step 1: Forward Pass: Compute model's predictions
        output = model(input_)

        # Step 2: Compute loss
        loss = loss_function(output, targets)

        # Correct prediction
        y_pred = torch.argmax(output, dim = 1)
        correct = torch.sum(y_pred == targets)

        batch_ct_train += 1

    # Step 3: Backward pass -Compute the gradients
    optimizer.zero_grad()

```

```

    loss.backward()

    # Gradient Clipping
    if grad_clipping:
        nn.utils.clip_grad_norm_(model.parameters(), max_norm=max_norm, norm_type=2)

    # Step 4: Update the parameters
    optimizer.step()

    # Add train loss of a batch
    running_train_loss += loss.item()

    # Add Corect counts of a batch
    running_train_correct += correct

    # log batch loss and accuracy
    if log_batch:
        if ((batch_ct_train + 1) % log_interval) == 0:
            wandb.log({"Train Batch Loss  ": loss})
            wandb.log({"Train Batch Acc  ": correct/len(targets)})

    # Calculate mean train loss for the whole dataset for a particular epoch
    train_loss = running_train_loss/len(train_loader)

    # Calculate accuracy for the whole dataset for a particular epoch
    train_acc = running_train_correct/len(train_loader.dataset)

    return train_loss, train_acc

def validate(valid_loader, loss_function, model, log_batch, log_interval):

    # initilalize variables as global
    # these counts will be updated every epoch
    global batch_ct_valid

    # Validation/Test loop
    # Initialize valid_loss at the he strat of the epoch
    running_val_loss = 0
    running_val_correct = 0

    # put the model in evaluation mode
    model.eval()

    with torch.no_grad():
        for input_,targets in valid_loader:

            # move inputs and outputs to GPUs
            input_ = input_.to(device)
            targets = targets.to(device)

            # Step 1: Forward Pass: Compute model's predictions
            output = model(input_)

```



```

# Step 2: Compute loss
loss = loss_function(output, targets)

# Correct Predictions
y_pred = torch.argmax(output, dim = 1)
correct = torch.sum(y_pred == targets)

batch_ct_valid += 1

# Add val loss of a batch
running_val_loss += loss.item()

# Add correct count for each batch
running_val_correct += correct

# log batch loss and accuracy
if log_batch:
    if ((batch_ct_valid + 1) % log_interval) == 0:
        wandb.log({f"Valid Batch Loss   ": loss})
        wandb.log({f"Valid Batch Accuracy ": correct/len(targets)})

# Calculate mean val loss for the whole dataset for a particular epoch
val_loss = running_val_loss/len(valid_loader)

# Calculate accuracy for the whole dataset for a particular epoch
val_acc = running_val_correct/len(valid_loader.dataset)

# scheduler step
# scheduler.step(valid_loss)
# scheduler.step()

return val_loss, val_acc

def train_loop(train_loader, valid_loader, model, optimizer, loss_function, epochs, device
               file_model):

    """
    Function for training the model and plotting the graph for train & validation loss vs ep
    Input: iterator for train dataset, initial weights and bias, epochs, learning rate, batch_size
    Output: final weights, bias and train loss and validation loss for each epoch.
    """

    # Create lists to store train and val loss at each epoch
    train_loss_history = []
    valid_loss_history = []
    train_acc_history = []
    valid_acc_history = []

    # initialize variables for early stopping

    delta = 0
    best_score = None
    valid_loss_min = np.Inf

```

```

counter_early_stop=0
early_stop=False

# Iterate for the given number of epochs
# Step 5: Repeat steps 1 - 4

for epoch in range(epochs):

    t0 = datetime.now()

    # Get train loss and accuracy for one epoch
    train_loss, train_acc = train(train_loader, loss_function, model, optimizer,
                                   wandb.config.grad_clipping, wandb.config.max_norm,
                                   wandb.config.log_batch, wandb.config.log_interval)
    valid_loss, valid_acc = validate(valid_loader, loss_function, model, wandb.config.lo

    dt = datetime.now() - t0

    # Save history of the Losses and accuracy
    train_loss_history.append(train_loss)
    train_acc_history.append(train_acc)

    valid_loss_history.append(valid_loss)
    valid_acc_history.append(valid_acc)

    # Log the train and valid loss to wandb
    wandb.log({"Train Loss ": train_loss, "epoch": epoch})
    wandb.log({"Train Acc ": train_acc, "epoch": epoch})

    wandb.log({"Valid Loss ": valid_loss, "epoch": epoch})
    wandb.log({"Valid Acc ": valid_acc, "epoch": epoch})

    if early_stopping:
        score = -valid_loss
        if best_score is None:
            best_score=score
            print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).')
            torch.save(model.state_dict(), file_model)
            valid_loss_min = valid_loss

        elif score < best_score + delta:
            counter_early_stop += 1
            print(f'Early stoping counter: {counter_early_stop} out of {patience}')
            if counter_early_stop > patience:
                early_stop = True

    else:
        best_score = score
        print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).')
        torch.save(model.state_dict(), file_model)
        counter_early_stop=0
        valid_loss_min = valid_loss

    if early_stop:

```

```

        print('Early Stopping')
        break

    else:

        score = -valid_loss
        if best_score is None:
            best_score=score
            print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).
            torch.save(model.state_dict(), file_model)
            valid_loss_min = valid_loss

        elif score < best_score + delta:
            print(f'Validation loss has not decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).

        else:
            best_score = score
            print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).
            torch.save(model.state_dict(), file_model)
            valid_loss_min = valid_loss

    # Print the train loss and accuracy for given number of epochs, batch size and number
    print(f'Epoch : {epoch+1} / {epochs}')
    print(f'Time to complete {epoch+1} is {dt}')
    # print(f'Learning rate: {scheduler._last_lr[0]}')
    print(f'Train Loss: {train_loss : .4f} | Train Accuracy: {train_acc * 100 : .4f}%')
    print(f'Valid Loss: {valid_loss : .4f} | Valid Accuracy: {valid_acc * 100 : .4f}%')
    print()
    torch.cuda.empty_cache()

    return train_loss_history, train_acc_history, valid_loss_history, valid_acc_history

def get_acc_pred(data_loader, model, device):

    """
    Function to get predictions and accuracy for a given data using estimated model
    Input: Data iterator, Final estimated weoights, bias
    Output: Prections and Accuracy for given dataset
    """

    # Array to store predicted labels
    predictions = torch.Tensor() # empty tensor
    predictions = predictions.to(device) # move predictions to GPU

    # Array to store actual labels
    y = torch.Tensor() # empty tensor
    y = y.to(device)

    # put the model in evaluation mode
    model.eval()

    # Iterate over batches from data iterator
    with torch.no_grad():
        for input_, targets in data_loader:

```

```

# move inputs and outputs to GPUs

input_ = input_.to(device)
targets = targets.to(device)

# Calculated the predicted labels
output = model(input_)

# Choose the label with maximum probability
prediction = torch.argmax(output, dim = 1)

# Add the predicted labels to the array
predictions = torch.cat((predictions, prediction))

# Add the actual labels to the array
y = torch.cat((y, targets))

# Check for complete dataset if actual and predicted labels are same or not
# Calculate accuracy
acc = (predictions == y).float().mean()

# Return tuple containing predictions and accuracy
return predictions, acc

```

## ▼ Logistic Regression

```

class LogisticRegression(torch.nn.Module):
    def __init__(self, input_dim, output_dim):
        super(LogisticRegression, self).__init__()
        self.linear = torch.nn.Linear(input_dim, output_dim)
    def forward(self, x):
        outputs = torch.sigmoid(self.linear(x))
        return outputs

epochs = 200
input_dim = 3*32*32 # Two inputs x1 and x2
output_dim = 1 # Single binary output
learning_rate = 0.01

model = LogisticRegression(input_dim,output_dim)

criterion = torch.nn.BCELoss()

optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

```

## ▼ One Hidden Layer

```

hyperparameters1 = SimpleNamespace(
    epochs = 20,
    input_dim = 28*28,
    output_dim = 10,
    h_sizes = [100] , # 1 Hidden Layer of 100 Neurons
    dprob = [0],
    non_linearity = nn.ReLU(),
    batch_norm = False,
    batch_size=25,
    learning_rate=0.07,
    dataset="FASHION_MNIST",
    architecture="MLP",
    log_interval = 1,
    log_batch = True,
    file_model = model_folder/'exp1_overfit_mnist.pt',
    grad_clipping = False, # DO NOT CHANGE hyperparameters below this
    early_stopping = False,
    max_norm = 1,
    momentum = 0,
    patience = 3,
    # scheduler_factor = 0.5,
    # scheduler_patience = 0,
    weight_decay = 0.00
)

```

```

wandb.config = hyperparameters1
wandb.config

```

```

namespace(architecture='MLP', batch_norm=False, batch_size=25,
dataset='FASHION_MNIST', dprob=[0], early_stopping=False, epochs=20,
file_model=PosixPath('/content/drive/MyDrive/Deep_Learning_UTD/Model/exp1_overfit_mni
grad_clipping=False, h_sizes=[100], input_dim=784, learning_rate=0.07,
log_batch=True, log_interval=1, max_norm=1, momentum=0, non_linearity=ReLU(),
output_dim=10, patience=3, weight_decay=0.0)

```

```

# Fix seed value
SEED = 2344
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

```

```

# Data Loader
train_loader = torch.utils.data.DataLoader(trainset, batch_size=wandb.config.batch_size, s
valid_loader = torch.utils.data.DataLoader(validset, batch_size=wandb.config.batch_size, s
# test_loader = torch.utils.data.DataLoader(testset, batch_size=wandb.config.batch_size,

```

```

# device
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

```

```

wandb.config.device = device

model = FashionMNIST_NN(wandb.config.input_dim, wandb.config.output_dim, wandb.config.h_si
                        wandb.config.dprob, wandb.config.non_linearity, wandb.config.bat

# Initialize weights from normal distribution with mean 0 and standard deviation 0.01
def init_weights(layer):
    if type(layer) == nn.Linear:
        torch.nn.init.kaiming_normal_(layer.weight, mean = 0, std = 0.01)
        # torch.nn.init.normal_(layer.weight, mean = 0, std = 0.001)
        torch.nn.init.zeros_(layer.bias)

model.to(wandb.config.device)
# model.apply(init_weights)

# loss_function
loss_function = nn.CrossEntropyLoss()

optimizer = torch.optim.Adam(model.parameters(),
                              lr = wandb.config.learning_rate,
                              weight_decay = wandb.config.weight_decay)

# scheduler = ReduceLROnPlateau(optimizer, mode='min', factor= wandb.config.scheduler_fact
#                               patience=wandb.config.scheduler_patience, verbose=True)

#scheduler = StepLR(optimizer, gamma=0.4,step_size=1, verbose=True)

# See live graphs in the notebook.
#%%wandb
batch_ct_train, batch_ct_valid = 0, 0
train_loss_history, train_acc_history, valid_loss_history, valid_acc_history = train_loop(

```

Validation loss has decreased (inf --> 2.400696). Saving Model...

Epoch : 1 / 20

Time to complete 1 is 0:00:13.682159

Train Loss: 2.6776 | Train Accuracy: 23.7604%

Valid Loss: 2.4007 | Valid Accuracy: 11.1750%

Validation loss has decreased (2.400696 --> 2.399744). Saving model...

Epoch : 2 / 20

Time to complete 2 is 0:00:13.766537

Train Loss: 2.8962 | Train Accuracy: 12.6042%

Valid Loss: 2.3997 | Valid Accuracy: 11.3500%

Validation loss has decreased (2.399744 --> 2.296860). Saving model...

```
Epoch : 3 / 20
Time to complete 3 is 0:00:13.968270
Train Loss: 2.6764 | Train Accuracy: 12.2667%
Valid Loss: 2.2969 | Valid Accuracy: 10.8417%
```

Validation loss has decreased (2.296860 --> 2.292094). Saving model...

```
Epoch : 4 / 20
Time to complete 4 is 0:00:13.827049
Train Loss: 2.3503 | Train Accuracy: 10.7208%
Valid Loss: 2.2921 | Valid Accuracy: 11.1250%
```

Validation loss has not decreased (2.292094 --> 2.376915). Not Saving Model...

```
Epoch : 5 / 20
Time to complete 5 is 0:00:16.397947
Train Loss: 2.3709 | Train Accuracy: 10.6208%
Valid Loss: 2.3769 | Valid Accuracy: 11.3500%
```

Validation loss has not decreased (2.292094 --> 2.317251). Not Saving Model...

```
Epoch : 6 / 20
Time to complete 6 is 0:00:14.738176
Train Loss: 2.3280 | Train Accuracy: 11.2979%
Valid Loss: 2.3173 | Valid Accuracy: 10.5917%
```

Validation loss has decreased (2.292094 --> 2.282408). Saving model...

```
Epoch : 7 / 20
Time to complete 7 is 0:00:14.257567
Train Loss: 2.3039 | Train Accuracy: 11.6854%
Valid Loss: 2.2824 | Valid Accuracy: 11.3500%
```

Validation loss has not decreased (2.282408 --> 2.301172). Not Saving Model...

```
Epoch : 8 / 20
Time to complete 8 is 0:00:14.579827
Train Loss: 2.2998 | Train Accuracy: 10.9000%
Valid Loss: 2.3012 | Valid Accuracy: 10.6000%
```

Validation loss has not decreased (2.282408 --> 2.311997). Not Saving Model...

```
Epoch : 9 / 20
Time to complete 9 is 0:00:14.404734
Train Loss: 2.2988 | Train Accuracy: 10.3146%
Valid Loss: 2.3120 | Valid Accuracy: 10.7417%
```

Validation loss has not decreased (2.282408 --> 2.321152). Not Saving Model...

```
Epoch : 10 / 20
Time to complete 10 is 0:00:15.625823
Train Loss: 2.2987 | Train Accuracy: 10.6187%
```

```
hyperparameters2 = SimpleNamespace(
    epochs = 20,
    input_dim = 28*28,
    output_dim = 10,
    h_sizes = [100,50] , # 2 Hidden Layer of 100,50 Neurons
    dprob = [0]*2,
    non_linearity = nn.ReLU(),
    batch_norm = False,
    batch_size=25,
    learning_rate=0.07,
    dataset="FASHION_MNIST",
    architecture="MLP",
```

```
log_interval = 1,
log_batch = True,
file_model = model_folder/'exp1_overfit_mnist.pt',
grad_clipping = False, # DO NOT CHANGE hyperparameters below this
early_stopping = False,
max_norm = 1,
momentum = 0,
patience = 3,
# scheduler_factor = 0.5,
# scheduler_patience = 0,
weight_decay = 0.00
)

wandb.config = hyperparameters2
wandb.config

# Fix seed value
SEED = 2344
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

# Data Loader
train_loader = torch.utils.data.DataLoader(trainset, batch_size=wandb.config.batch_size, s
valid_loader = torch.utils.data.DataLoader(validset, batch_size=wandb.config.batch_size, s
# test_loader = torch.utils.data.DataLoader(testset, batch_size=wandb.config.batch_size,

# device
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

wandb.config.device = device

model2 = FashionMNIST_NN(wandb.config.input_dim, wandb.config.output_dim, wandb.config.h_s
                        wandb.config.dprob, wandb.config.non_linearity, wandb.config.bat

# Initialize weights from normal distribution with mean 0 and standard deviation 0.01
def init_weights(layer):
    if type(layer) == nn.Linear:
        torch.nn.init.kaiming_normal_(layer.weight, mean = 0, std = 0.01)
        # torch.nn.init.normal_(layer.weight, mean = 0, std = 0.001)
        torch.nn.init.zeros_(layer.bias)

model2.to(wandb.config.device)
# model.apply(init_weights)

# loss_function
loss_function = nn.CrossEntropyLoss()

optimizer = torch.optim.Adam(model2.parameters(),
                              lr = wandb.config.learning_rate,
                              weight_decay = wandb.config.weight_decay)
```



```
# scheduler = ReduceLROnPlateau(optimizer, mode='min', factor= wandb.config.scheduler_fact
#                               patience=wandb.config.scheduler_patience, verbose=True)

#scheduler = StepLR(optimizer, gamma=0.4,step_size=1, verbose=True)

# See live graphs in the notebook.
#%%wandb
batch_ct_train, batch_ct_valid = 0, 0
train_loss_history, train_acc_history, valid_loss_history, valid_acc_history = train_loop(
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 4m 49s completed at 4:52 PM

