

▼ Load Libraries/Install Software

```
if 'google.colab' in str(get_ipython()):
    print('Running on Colab')
else:
    print('Not running on Colab')
```

Running on Colab

```
if 'google.colab' in str(get_ipython()):
    !pip install wandb --upgrade -q
```

1.9 MB	27.3 MB/s
182 kB	17.7 MB/s
166 kB	49.7 MB/s
63 kB	1.8 MB/s
166 kB	49.4 MB/s
162 kB	72.0 MB/s
162 kB	72.4 MB/s
158 kB	71.6 MB/s
157 kB	72.9 MB/s
157 kB	68.6 MB/s
157 kB	70.8 MB/s
157 kB	70.4 MB/s
157 kB	72.4 MB/s
157 kB	69.1 MB/s
157 kB	75.2 MB/s
156 kB	71.1 MB/s

Building wheel for pathtools (setup.py) ... done

```
if 'google.colab' in str(get_ipython()):
    from google.colab import drive
    drive.mount('/content/drive')
```

Mounted at /content/drive

#Importing the libraries

```
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F
```

```
from torch.optim.lr_scheduler import ReduceLROnPlateau, OneCycleLR, CyclicLR, ExponentialLR
```

```
import numpy as np
import random
```

```
from datetime import datetime
```

```
from pathlib import Path
import sys
from types import SimpleNamespace
```

```
import wandb
```

```
#Login into wandb
```

```
wandb.login()
```

```
ERROR:wandb.jupyter:Failed to detect the name of this notebook, you can set it manually
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
True
```

```
wandb.init(name = "Hw6_CIFAR10.ipynb", project = 'dl22_HW6')
```

```
wandb: Currently logged in as: pranavshekhar2. Use `wandb login --relogin` to force relogin
Tracking run with wandb version 0.13.5
Run data is saved locally in /content/wandb/run-20221112_005737-fn8drw2v
Syncing run Hw6_CIFAR10.ipynb to Weights & Biases (docs)
```

▼ Specify Project Folders

```
# This is the path where we will download and save data
if 'google.colab' in str(get_ipython()):
    data_folder = Path('/content/drive/MyDrive/Deep_Learning_UTD/Dataset')
    model_folder = Path('/content/drive/MyDrive/Deep_Learning_UTD/Model')
else:
    base_folder = Path('/home/harpreet/Insync/google_drive_shaanmoor/data')
```

▼ Load CIFAR10 Dataset

```
trans1 = transforms.ToTensor()
trans2 = transforms.Normalize((0.49, 0.482, 0.447), (0.247, 0.244, 0.262))

trans = transforms.Compose([trans1, trans2])

#Train and validation set

train_val_set = torchvision.datasets.CIFAR10(root = data_folder,
                                              transform = trans,
                                              train = True,
                                              download = True)

testset = torchvision.datasets.CIFAR10(root = data_folder,
                                         transform = trans,
```

```

train = False,
download = True
)

```

```

Files already downloaded and verified
Files already downloaded and verified

```

```

def split_dataset(base_dataset, fraction, seed):
    split_a_size = int(fraction * len(base_dataset))
    split_b_size = len(base_dataset) - split_a_size
    return torch.utils.data.random_split(base_dataset, [split_a_size, split_b_size], gener
)

```

```

trainset, validset = split_dataset(train_val_set,0.8,10)

```

```

# Since the transforms are not applied, we will manually first divide by 255
# we will then get the mean and std dev
# the images are still in mumpy with the shape (number of images, H, W, Channels)
# Since we need mean, std dev
train_data = train_val_set.data[trainset.indices]/255
train_data.shape
print(train_data.mean(axis = (0,1,2)))
print(train_data.std(axis = (0,1,2)))

```

```

[0.49114078 0.48191055 0.44641415]
[0.24707852 0.24355821 0.26162066]

```

▼ Checking inputs for CIFAR10 Dataset

```

print(type(trainset), type(trainset.dataset), type(trainset.indices), sep = '\n')

```

```

<class 'torch.utils.data.dataset.Subset'>
<class 'torchvision.datasets.cifar.CIFAR10'>
<class 'list'>

```

```

trainset.indices[0:5]

```

```

[46937, 45069, 32498, 25031, 16172]

```

```

print(len(trainset), len(trainset.dataset), len(trainset.indices), sep = '\n')

```

```

40000
50000
40000

```

```

# Shape of training data
len(trainset.indices), len(validset.indices)

```

```
(40000, 10000)

print(train_val_set.data[trainset.indices].shape)
print(trainset.dataset.data[trainset.indices].shape)

(40000, 32, 32, 3)
(40000, 32, 32, 3)

train_val_set.data[validset.indices].shape

(10000, 32, 32, 3)

# Shape of testing data
testset.data.shape

(10000, 32, 32, 3)

# check the max value of inputs - the transformation are not yet applied.
# the transofrmation are applied iteratively on batches
# when we craete batch by iterating over dataloader
train_val_set.data[trainset.indices].max()

255

# check the min value of inputs
train_val_set.data[trainset.indices].min()

0
```

▼ Visualize the Data


```
# Initializing the batch size
batch_size = 32

check_loader = torch.utils.data.DataLoader(trainset, batch_size = 64, shuffle = True)

# Get some random training images
dataiter = iter(check_loader)
images, labels = dataiter.next()

# Create grid of images
img_grid = torchvision.utils.make_grid(images[0:50], nrow = 10)

# Logging to W&B
images = wandb.Image(img_grid, caption = "Sample images")
images.image
```



▼ 1. Implement CNN on CIFAR10 Dataset

```
from torch.nn.modules import Conv2d
class Cifar10CNN(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1_layer = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=5, padding=2, bias=True),
            nn.ReLU(),
            nn.Conv2d(in_channels=32, out_channels=16, kernel_size=
            nn.ReLU(),

        )

        self.flatten = nn.Flatten()
        self.fc2 = nn.Linear(16*32*32, out_features=10)

    def forward(self,x):
        x = self.conv1_layer(x)
        x = self.flatten(x)
        x = F.relu(self.fc2(x))

        return x
```

```
from torchsummary import summary
summary(Cifar10CNN().cuda(), (3, 32, 32))
```

```
-----
Layer (type)          Output Shape          Param #
=====
Conv2d-1              [-1, 32, 32, 32]      2,432
ReLU-2                [-1, 32, 32, 32]      0
Conv2d-3              [-1, 16, 32, 32]      4,624
ReLU-4                [-1, 16, 32, 32]      0
Flatten-5             [-1, 16384]           0
Linear-6               [-1, 10]              163,850
=====
Total params: 170,906
Trainable params: 170,906
Non-trainable params: 0
-----
Input size (MB): 0.01
```

Forward/backward pass size (MB): 0.88

Params size (MB): 0.65

Estimated Total Size (MB): 1.54

▼ Function for Training Loop

```
def train(train_loader, loss_function, model, optimizer, grad_clipping, max_norm, log_batch_size):
    # Training Loop

    # initialize variables as global
    # these counts will be updated every epoch
    global batch_ct_train

    # Initialize train_loss at the he start of the epoch
    running_train_loss = 0
    running_train_correct = 0

    # put the model in training mode
    model.train()
    # Iterate on batches from the dataset using train_loader
    for input_, targets in train_loader:

        # move inputs and outputs to GPUs
        input_ = input_.to(device)
        targets = targets.to(device)

        # Step 1: Forward Pass: Compute model's predictions
        output = model(input_)

        # Step 2: Compute loss
        loss = loss_function(output, targets)

        # Correct prediction
        y_pred = torch.argmax(output, dim = 1)
        correct = torch.sum(y_pred == targets)

        batch_ct_train += 1

        # Step 3: Backward pass -Compute the gradients
        optimizer.zero_grad()
        loss.backward()

        # Gradient Clipping
        if grad_clipping:
            nn.utils.clip_grad_norm_(model.parameters(), max_norm=max_norm, norm_type=2)

        # Step 4: Update the parameters
        optimizer.step()
```

```

# For One Cycle Learning Rate - add step here
# scheduler.step()

# Add train loss of a batch
running_train_loss += loss.item()

# Add Corect counts of a batch
running_train_correct += correct

# log batch loss and accuracy
if log_batch:
    if ((batch_ct_train + 1) % log_interval) == 0:
        wandb.log({f"Train Batch Loss  ": loss})
        wandb.log({f"Train Batch Acc  ": correct/len(targets)})
        # print(f'Learning rate: {scheduler.get_last_lr()}')

# Calculate mean train loss for the whole dataset for a particular epoch
train_loss = running_train_loss/len(train_loader)

# Calculate accuracy for the whole dataset for a particular epoch
train_acc = running_train_correct/len(train_loader.dataset)

return train_loss, train_acc

```

▼ Function for Validation Loops

```

def validate(valid_loader, loss_function, model, log_batch, log_interval):

    # inititalize variables as global
    # these counts will be updated every epoch
    global batch_ct_valid

    # Validation/Test loop
    # Initialize valid_loss at the he strat of the epoch
    running_val_loss = 0
    running_val_correct = 0

    # put the model in evaluation mode
    model.eval()

    with torch.no_grad():
        for input_, targets in valid_loader:

            # move inputs and outputs to GPUs
            input_ = input_.to(device)
            targets = targets.to(device)

            # Step 1: Forward Pass: Compute model's predictions
            output = model(input_)

```

```

# Step 2: Compute loss
loss = loss_function(output, targets)

# Correct Predictions
y_pred = torch.argmax(output, dim = 1)
correct = torch.sum(y_pred == targets)

batch_ct_valid += 1

# Add val loss of a batch
running_val_loss += loss.item()

# Add correct count for each batch
running_val_correct += correct

# log batch loss and accuracy
if log_batch:
    if ((batch_ct_valid + 1) % log_interval) == 0:
        wandb.log({f"Valid Batch Loss   ": loss})
        wandb.log({f"Valid Batch Accuracy ": correct/len(targets)})

# Calculate mean val loss for the whole dataset for a particular epoch
val_loss = running_val_loss/len(valid_loader)

# Calculate accuracy for the whole dataset for a particular epoch
val_acc = running_val_correct/len(valid_loader.dataset)

# scheduler step
# scheduler.step(val_loss)
# scheduler.step()

return val_loss, val_acc

```

▼ Function for Model Training

```

def train_loop(train_loader, valid_loader, model, optimizer, loss_function, epochs, device,
               file_model, save_best_model):

    """
    Function for training the model and plotting the graph for train & validation loss vs ep
    Input: iterator for train dataset, initial weights and bias, epochs, learning rate, batch_size
    Output: final weights, bias and train loss and validation loss for each epoch.
    """

    # Create lists to store train and val loss at each epoch
    train_loss_history = []
    valid_loss_history = []
    train_acc_history = []
    valid_acc_history = []

    # initialize variables for early stopping

```



```

delta = 0
best_score = None
valid_loss_min = np.Inf
counter_early_stop=0
early_stop=False

# Iterate for the given number of epochs
# Step 5: Repeat steps 1 - 4

for epoch in range(epochs):

    t0 = datetime.now()

    # Get train loss and accuracy for one epoch
    train_loss, train_acc = train(train_loader, loss_function, model, optimizer,
                                   wandb.config.grad_clipping, wandb.config.max_norm,
                                   wandb.config.log_batch, wandb.config.log_interval)
    valid_loss, valid_acc = validate(valid_loader, loss_function, model, wandb.config.lo

    dt = datetime.now() - t0

    # Save history of the Losses and accuracy
    train_loss_history.append(train_loss)
    train_acc_history.append(train_acc)

    valid_loss_history.append(valid_loss)
    valid_acc_history.append(valid_acc)

    # Log the train and valid loss to wandb
    wandb.log({"Train Loss ": train_loss, "epoch": epoch})
    wandb.log({"Train Acc ": train_acc, "epoch": epoch})

    wandb.log({"Valid Loss ": valid_loss, "epoch": epoch})
    wandb.log({"Valid Acc ": valid_acc, "epoch": epoch})

    if early_stopping:
        score = -valid_loss
        if best_score is None:
            best_score=score
            print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).
            torch.save(model.state_dict(), file_model)
            valid_loss_min = valid_loss

        elif score < best_score + delta:
            counter_early_stop += 1
            print(f'Early stoping counter: {counter_early_stop} out of {patience}')
            if counter_early_stop > patience:
                early_stop = True

    else:
        best_score = score
        print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).
        torch.save(model.state_dict(), file_model)

```

```

        counter_early_stop=0
        valid_loss_min = valid_loss

    if early_stop:
        print('Early Stopping')
        break

    elif save_best_model:

        score = -valid_loss
        if best_score is None:
            best_score=score
            print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).')
            torch.save(model.state_dict(), file_model)
            valid_loss_min = valid_loss

        elif score < best_score + delta:
            print(f'Validation loss has not decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).')

        else:
            best_score = score
            print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).')
            torch.save(model.state_dict(), file_model)
            valid_loss_min = valid_loss

    else:
        torch.save(model.state_dict(), file_model)

    # Print the train loss and accuracy for given number of epochs, batch size and number
    print(f'Epoch : {epoch+1} / {epochs}')
    print(f'Time to complete {epoch+1} is {dt}')
    # print(f'Learning rate: {scheduler._last_lr[0]}') # for ReduceLR
    # print(f'Learning rate: {scheduler.get_last_lr()}') # for OneCycle
    print(f'Train Loss: {train_loss : .4f} | Train Accuracy: {train_acc * 100 : .4f}%')
    print(f'Valid Loss: {valid_loss : .4f} | Valid Accuracy: {valid_acc * 100 : .4f}%')
    print()
    torch.cuda.empty_cache()

    return train_loss_history, train_acc_history, valid_loss_history, valid_acc_history

```

▼ Function for Accuracy and Predictions

```

def get_acc_pred(data_loader, model, device):

    """
    Function to get predictions and accuracy for a given data using estimated model
    Input: Data iterator, Final estimated weoights, bias
    Output: Prections and Accuracy for given dataset
    """

    # Array to store predicted labels
    predictions = torch.Tensor() # empty tensor

```

```

predictions = predictions.to(device) # move predictions to GPU

# Array to store actual labels
y = torch.Tensor() # empty tensor
y = y.to(device)

# put the model in evaluation mode
model.eval()

# Iterate over batches from data iterator
with torch.no_grad():
    for input_, targets in data_loader:

        # move inputs and outputs to GPUs

        input_ = input_.to(device)
        targets = targets.to(device)

        # Calculated the predicted labels
        output = model(input_)

        # Choose the label with maximum probability
        prediction = torch.argmax(output, dim = 1)

        # Add the predicted labels to the array
        predictions = torch.cat((predictions, prediction))

        # Add the actual labels to the array
        y = torch.cat((y, targets))

# Check for complete dataset if actual and predicted labels are same or not
# Calculate accuracy
acc = (predictions == y).float().mean()

# Return tuple containing predictions and accuracy
return predictions, acc

```

▼ Meta Data

```

hyperparameters = SimpleNamespace(
    epochs = 20,
    output_dim = 10,
    batch_size=256,
    learning_rate=1e-2,
    dataset="Cifar10",
    architecture="Cifar10_CustomNN",
    log_interval = 100,
    log_batch = True,
    file_model = model_folder/'cifar10_custom_model.pt',
    grad_clipping = True,
    max_norm = 1,
    momentum = 0.9,

```

```

patience = 5,
early_stopping = True,
scheduler_factor = 0.5,
scheduler_patience = 0,
save_best_model = False,
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu'),
weight_decay = 0,

)

```

▼ Data Loaders, Loss Function, Optimizer

```

wandb.config = hyperparameters
wandb.config

```

```

namespace(architecture='Cifar10_CustomNN', batch_size=256, dataset='Cifar10',
device=device(type='cuda', index=0), early_stopping=True, epochs=20,
file_model=PosixPath('/content/drive/MyDrive/Deep_Learning_UTD/Model/cifar10_custom_n
grad_clipping=True, learning_rate=0.01, log_batch=True, log_interval=100,
max_norm=1, momentum=0.9, output_dim=10, patience=5, save_best_model=False,
scheduler_factor=0.5, scheduler_patience=0, weight_decay=0)

```

```

# Fix seed value
SEED = 2345
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

# Data Loader
train_loader = torch.utils.data.DataLoader(trainset, batch_size=wandb.config.batch_size, s
valid_loader = torch.utils.data.DataLoader(validset, batch_size=wandb.config.batch_size, s
test_loader = torch.utils.data.DataLoader(testset, batch_size=wandb.config.batch_size, s

# model
model = Cifar10CNN()

# Initialize weights from normal distribution with mean 0 and standard deviation 0.01
def init_weights(m):
    if type(m) == nn.Linear:
        torch.nn.init.kaiming_normal_(m.weight)
        torch.nn.init.zeros_(m.bias)

    if type(m) == nn.Conv2d:
        torch.nn.init.kaiming_normal_(m.weight)
        torch.nn.init.zeros_(m.bias)

model.to(wandb.config.device)

```

```
# model.apply(init_weights)

# loss_function
loss_function = nn.CrossEntropyLoss()

optimizer = torch.optim.SGD(model.parameters(),
                              lr = wandb.config.learning_rate,
                              momentum = wandb.config.momentum,
                              weight_decay = wandb.config.weight_decay)

# scheduler = ReduceLROnPlateau(optimizer, mode='min', factor= wandb.config.scheduler_fact
#                               patience=wandb.config.scheduler_patience, verbose=True)
```

▼ Training Model

```
wandb.watch(model, log = 'all', log_freq=25, log_graph=True)
```

```
wandb: logging graph, to disable use `wandb.watch(log_graph=False)`
[<wandb.wandb_torch.TorchGraph at 0x7f905ba45190>]
```

```
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
```

```
# See live graphs in the notebook.
```

```
##wandb
```

```
batch_ct_train, batch_ct_valid = 0, 0
```

```
train_loss_history, train_acc_history, valid_loss_history, valid_acc_history = train_loop(
```

```
Validation loss has decreased (inf --> 1.951506). Saving Model...
```

```
Epoch : 1 / 20
```

```
Time to complete 1 is 0:00:13.998869
```

```
Train Loss: 2.0739 | Train Accuracy: 29.4625%
```

```
Valid Loss: 1.9515 | Valid Accuracy: 34.0700%
```

```
Validation loss has decreased (1.951506 --> 1.702215). Saving model...
```

```
Epoch : 2 / 20
```

```
Time to complete 2 is 0:00:13.409972
```

```
Train Loss: 1.7974 | Train Accuracy: 39.9425%
```

```
Valid Loss: 1.7022 | Valid Accuracy: 42.1400%
```

```
Validation loss has decreased (1.702215 --> 1.645574). Saving model...
Epoch : 3 / 20
Time to complete 3 is 0:00:13.344534
Train Loss: 1.6065 | Train Accuracy: 45.0850%
Valid Loss: 1.6456 | Valid Accuracy: 43.4800%

Validation loss has decreased (1.645574 --> 1.612931). Saving model...
Epoch : 4 / 20
Time to complete 4 is 0:00:13.294269
Train Loss: 1.5368 | Train Accuracy: 47.4725%
Valid Loss: 1.6129 | Valid Accuracy: 44.1600%

Validation loss has decreased (1.612931 --> 1.566517). Saving model...
Epoch : 5 / 20
Time to complete 5 is 0:00:13.259495
Train Loss: 1.4920 | Train Accuracy: 49.0575%
Valid Loss: 1.5665 | Valid Accuracy: 46.7900%

Validation loss has decreased (1.566517 --> 1.514880). Saving model...
Epoch : 6 / 20
Time to complete 6 is 0:00:13.293307
Train Loss: 1.4455 | Train Accuracy: 50.6550%
Valid Loss: 1.5149 | Valid Accuracy: 48.2700%

Early stoping counter: 1 out of 5
Epoch : 7 / 20
Time to complete 7 is 0:00:14.512335
Train Loss: 1.4106 | Train Accuracy: 51.6950%
Valid Loss: 1.5471 | Valid Accuracy: 47.1500%

Validation loss has decreased (1.514880 --> 1.482010). Saving model...
Epoch : 8 / 20
Time to complete 8 is 0:00:13.132618
Train Loss: 1.3775 | Train Accuracy: 52.6375%
Valid Loss: 1.4820 | Valid Accuracy: 49.1100%

Early stoping counter: 1 out of 5
Epoch : 9 / 20
Time to complete 9 is 0:00:14.502721
Train Loss: 1.3442 | Train Accuracy: 53.7000%
Valid Loss: 1.4836 | Valid Accuracy: 49.1100%

Early stoping counter: 2 out of 5
Epoch : 10 / 20
Time to complete 10 is 0:00:14.450005
Train Loss: 1.3113 | Train Accuracy: 54.8675%
```

▼ 2. Experiment different CNN on CIFAR10 Dataset

```
#Subset of 50 images
```

```
# n sample points
train_sample_size = 50
```

```
# Getting n random indices
train_subset_indices = random.sample(range(0, len(trainset)), train_sample_size)
```

```
# Getting subset of dataset
train_subset = torch.utils.data.Subset(trainset, train_subset_indices)

train_subset.dataset.data[train_subset.indices].shape

(50, 32, 32, 3)
```

Experiment 1 - Denser Network and Wider Network

Conv (Kernel Size = 10)->ReLU->MaxPool (Kernel Size = 2, Stride = 2)

Conv (Kernel Size = 5)->ReLU->BatchNorm

Conv (Kernel Size = 3)->ReLU->MaxPool (Kernel Size = 2, Stride = 2)

Flatten-> Linear_Layer

Optimizer = SGD

LR Scheduler - ReduceLROnPlateau

Experiment 2 - Denser Network and Wider Network

Conv (Kernel Size = 12)->ELU->AvgPool (Kernel Size = 2, Stride = 2)

Conv (Kernel Size = 5)->ELU->BatchNorm

Conv (Kernel Size = 3)->ELU->AvgPool (Kernel Size = 2, Stride = 2)

Flatten-> Dropout_Layer-> Linear_Layer

Optimizer = SGD

LR Scheduler - OneCycleLR

Experiment 3 - Denser Network and Wider Network

Conv (Kernel Size = 12)->ReLU->MaxPool (Kernel Size = 2, Stride = 2)

Conv (Kernel Size = 3)->ReLU->BatchNorm

Conv (Kernel Size = 2)->ReLU->MaxPool (Kernel Size = 2, Stride = 2)

Flatten-> Linear_Layer

Optimizer = Adam

LR Scheduler - ReduceLROnPlateau

Experiment 4 - Denser Network and Wider Network

Conv (Kernel Size = 10) ->ReLU->MaxPool (Kernel Size = 2 , stride = 2)

Conv (Kernel Size = 5) ->ReLU->BatchNorm

Conv (Kernel Size = 3) ->ReLU->MaxPool (Kernel Size = 2 , stride = 2)

Flatten-> Linear_Layer

Optimizer = Adam

LR Scheduler - ReduceLROnPlateau

Experiment 5 - Denser Network and Wider Network

Conv (Kernel Size = 10) -> ReLU->AvgPool (Kernel Size = 2, Stride = 2)

Conv (Kernel Size = 5) -> ReLU->BatchNorm

Conv (Kernel Size = 3) -> ReLU->AvgPool (Kernel Size = 2, Stride = 2)

Flatten-> Linear_Layer

Optimizer = Adam

LR Scheduler - ReduceLROnPlateau

▼ Best Model

Conv (Kernel Size = 10)->ReLU->MaxPool (Kernel Size = 2, Stride = 2)

Conv (Kernel Size = 5)->ReLU->BatchNorm

Conv (Kernel Size = 3)->ReLU->MaxPool (Kernel Size = 2, Stride = 2)

Flatten-> Linear_Layer

Optimizer = SGD

LR Scheduler - ReduceLROnPlateau

```
from torch.nn.modules import Conv2d
class Cifar10CNN_Experiment(nn.Module):
```



```

def __init__(self):
    super().__init__()

    self.conv1_layer = nn.Sequential(
        nn.Conv2d(in_channels=3, out_channels=64, kernel_size=10, padding=2, bias=True),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size = 2, stride = 2),
        nn.Conv2d(in_channels=64, out_channels=256, kernel_size=5, padding=2, bias=True),
        nn.ReLU(),
        nn.BatchNorm2d(256,momentum=0.9),
        nn.Conv2d(in_channels=256, out_channels=16, kernel_size=3, padding=2, bias=True),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size = 2, stride = 2)
    )

    self.flatten = nn.Flatten()
    self.fc2 = nn.Linear(16*7*7, out_features=10)

def forward(self,x):
    x = self.conv1_layer(x)
    x = self.flatten(x)
    x = F.relu(self.fc2(x))

    return x

```

```

from torchsummary import summary
summary(Cifar10CNN_Experiment().cuda(), (3, 32, 32))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 27, 27]	19,264
ReLU-2	[-1, 64, 27, 27]	0
MaxPool2d-3	[-1, 64, 13, 13]	0
Conv2d-4	[-1, 256, 13, 13]	409,856
ReLU-5	[-1, 256, 13, 13]	0
BatchNorm2d-6	[-1, 256, 13, 13]	512
Conv2d-7	[-1, 16, 15, 15]	36,880
ReLU-8	[-1, 16, 15, 15]	0
MaxPool2d-9	[-1, 16, 7, 7]	0
Flatten-10	[-1, 784]	0
Linear-11	[-1, 10]	7,850

=====
Total params: 474,362

Trainable params: 474,362

Non-trainable params: 0

Input size (MB): 0.01

Forward/backward pass size (MB): 1.85

Params size (MB): 1.81

Estimated Total Size (MB): 3.67

▼ Function for Training Loop

```
def train(train_loader, loss_function, model, optimizer, grad_clipping, max_norm, log_batch
```

```

    # Training Loop

    # initialize variables as global
    # these counts will be updated every epoch
    global batch_ct_train

    # Initialize train_loss at the he start of the epoch
    running_train_loss = 0
    running_train_correct = 0

    # put the model in training mode

    model.train()
    # Iterate on batches from the dataset using train_loader
    for input_, targets in train_loader:

        # move inputs and outputs to GPUs
        input_ = input_.to(device)
        targets = targets.to(device)

        # Step 1: Forward Pass: Compute model's predictions
        output = model(input_)

        # Step 2: Compute loss
        loss = loss_function(output, targets)

        # Correct prediction
        y_pred = torch.argmax(output, dim = 1)
        correct = torch.sum(y_pred == targets)

        batch_ct_train += 1

        # Step 3: Backward pass -Compute the gradients
        optimizer.zero_grad()
        loss.backward()

        # Gradient Clipping
        if grad_clipping:
            nn.utils.clip_grad_norm_(model.parameters(), max_norm=max_norm, norm_type=2)

        # Step 4: Update the parameters
        optimizer.step()

        # For One Cycle Learning Rate - add step here
        # scheduler.step()

        # Add train loss of a batch
        running_train_loss += loss.item()

```

```

# Add Corect counts of a batch
running_train_correct += correct

# log batch loss and accuracy
if log_batch:
    if ((batch_ct_train + 1) % log_interval) == 0:
        wandb.log({"Train Batch Loss  ": loss})
        wandb.log({"Train Batch Acc  ": correct/len(targets)})
        # print(f'Learning rate: {scheduler.get_last_lr()}')

# Calculate mean train loss for the whole dataset for a particular epoch
train_loss = running_train_loss/len(train_loader)

# Calculate accuracy for the whole dataset for a particular epoch
train_acc = running_train_correct/len(train_loader.dataset)

return train_loss, train_acc

```

▼ Function for Validation Loops

```

def validate(valid_loader, loss_function, model, log_batch, log_interval):

    # initilalize variables as global
    # these counts will be updated every epoch
    global batch_ct_valid

    # Validation/Test loop
    # Initialize valid_loss at the he strat of the epoch
    running_val_loss = 0
    running_val_correct = 0

    # put the model in evaluation mode
    model.eval()

    with torch.no_grad():
        for input_,targets in valid_loader:

            # move inputs and outputs to GPUs
            input_ = input_.to(device)
            targets = targets.to(device)

            # Step 1: Forward Pass: Compute model's predictions
            output = model(input_)

            # Step 2: Compute loss
            loss = loss_function(output, targets)

            # Correct Predictions
            y_pred = torch.argmax(output, dim = 1)
            correct = torch.sum(y_pred == targets)

```

```

    batch_ct_valid += 1

    # Add val loss of a batch
    running_val_loss += loss.item()

    # Add correct count for each batch
    running_val_correct += correct

    # log batch loss and accuracy
    if log_batch:
        if ((batch_ct_valid + 1) % log_interval) == 0:
            wandb.log({"Valid Batch Loss  ": loss})
            wandb.log({"Valid Batch Accuracy ": correct/len(targets)})

    # Calculate mean val loss for the whole dataset for a particular epoch
    val_loss = running_val_loss/len(valid_loader)

    # Calculate accuracy for the whole dataset for a particular epoch
    val_acc = running_val_correct/len(valid_loader.dataset)

    # scheduler step
    scheduler.step(val_loss)

return val_loss, val_acc

```

▼ Function for Model Training

```

def train_loop(train_loader, valid_loader, model, optimizer, loss_function, epochs, device,
               file_model, save_best_model):

    """
    Function for training the model and plotting the graph for train & validation loss vs ep
    Input: iterator for train dataset, initial weights and bias, epochs, learning rate, batch_size
    Output: final weights, bias and train loss and validation loss for each epoch.
    """

    # Create lists to store train and val loss at each epoch
    train_loss_history = []
    valid_loss_history = []
    train_acc_history = []
    valid_acc_history = []

    # initialize variables for early stopping

    delta = 0
    best_score = None
    valid_loss_min = np.Inf
    counter_early_stop=0
    early_stop=False

    # Iterate for the given number of epochs

```

```
# Step 5: Repeat steps 1 - 4
```

```
for epoch in range(epochs):
```

```
    t0 = datetime.now()
```

```
    # Get train loss and accuracy for one epoch
```

```
    train_loss, train_acc = train(train_loader, loss_function, model, optimizer,
                                   wandb.config.grad_clipping, wandb.config.max_norm,
                                   wandb.config.log_batch, wandb.config.log_interval)
```

```
    valid_loss, valid_acc = validate(valid_loader, loss_function, model, wandb.config.lo
```

```
    dt = datetime.now() - t0
```

```
    # Save history of the Losses and accuracy
```

```
    train_loss_history.append(train_loss)
```

```
    train_acc_history.append(train_acc)
```

```
    valid_loss_history.append(valid_loss)
```

```
    valid_acc_history.append(valid_acc)
```

```
    # Log the train and valid loss to wandb
```

```
    wandb.log({"Train Loss ": train_loss, "epoch": epoch})
```

```
    wandb.log({"Train Acc ": train_acc, "epoch": epoch})
```

```
    wandb.log({"Valid Loss ": valid_loss, "epoch": epoch})
```

```
    wandb.log({"Valid Acc ": valid_acc, "epoch": epoch})
```

```
    if early_stopping:
```

```
        score = -valid_loss
```

```
        if best_score is None:
```

```
            best_score=score
```

```
            print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).
```

```
            torch.save(model.state_dict(), file_model)
```

```
            valid_loss_min = valid_loss
```

```
        elif score < best_score + delta:
```

```
            counter_early_stop += 1
```

```
            print(f'Early stoping counter: {counter_early_stop} out of {patience}')
```

```
            if counter_early_stop > patience:
```

```
                early_stop = True
```

```
        else:
```

```
            best_score = score
```

```
            print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).
```

```
            torch.save(model.state_dict(), file_model)
```

```
            counter_early_stop=0
```

```
            valid_loss_min = valid_loss
```

```
    if early_stop:
```

```
        print('Early Stopping')
```

```
        break
```

```
    elif save_best_model:
```

```

score = -valid_loss
if best_score is None:
    best_score=score
    print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).
    torch.save(model.state_dict(), file_model)
    valid_loss_min = valid_loss

elif score < best_score + delta:
    print(f'Validation loss has not decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).

else:
    best_score = score
    print(f'Validation loss has decreased ({valid_loss_min:.6f} --> {valid_loss:.6f}).
    torch.save(model.state_dict(), file_model)
    valid_loss_min = valid_loss

else:
    torch.save(model.state_dict(), file_model)

# Print the train loss and accuracy for given number of epochs, batch size and number
print(f'Epoch : {epoch+1} / {epochs}')
print(f'Time to complete {epoch+1} is {dt}')
print(f'Learning rate: {scheduler._last_lr[0]}') # for ReduceLR
# print(f'Learning rate: {scheduler.get_last_lr()}') # for OneCycle
print(f'Train Loss: {train_loss : .4f} | Train Accuracy: {train_acc * 100 : .4f}%')
print(f'Valid Loss: {valid_loss : .4f} | Valid Accuracy: {valid_acc * 100 : .4f}%')
print()
torch.cuda.empty_cache()

return train_loss_history, train_acc_history, valid_loss_history, valid_acc_history

```

▼ Function for Accuracy and Predictions

```

def get_acc_pred(data_loader, model, device):

    """
    Function to get predictions and accuracy for a given data using estimated model
    Input: Data iterator, Final estimated weights, bias
    Output: Predictions and Accuracy for given dataset
    """

    # Array to store predicted labels
    predictions = torch.Tensor() # empty tensor
    predictions = predictions.to(device) # move predictions to GPU

    # Array to store actual labels
    y = torch.Tensor() # empty tensor
    y = y.to(device)

    # put the model in evaluation mode
    model.eval()

```

```
# Iterate over batches from data iterator
with torch.no_grad():
    for input_, targets in data_loader:

        # move inputs and outputs to GPUs

        input_ = input_.to(device)
        targets = targets.to(device)

        # Calculated the predicted labels
        output = model(input_)

        # Choose the label with maximum probability
        prediction = torch.argmax(output, dim = 1)

        # Add the predicted labels to the array
        predictions = torch.cat((predictions, prediction))

        # Add the actual labels to the array
        y = torch.cat((y, targets))

# Check for complete dataset if actual and predicted labels are same or not
# Calculate accuracy
acc = (predictions == y).float().mean()

# Return tuple containing predictions and accuracy
return predictions, acc
```

▼ Meta Data

```
hyperparameters2 = SimpleNamespace(
    epochs = 10,
    output_dim = 10,
    batch_size=256,
    learning_rate=1e-2,
    dataset="Cifar10",
    architecture="Cifar10_CustomNN_Experiment",
    log_interval = 100,
    log_batch = True,
    file_model = model_folder/'cifar10_custom_model_experiment.pt',
    grad_clipping = True,
    max_norm = 1,
    momentum = 0.9,
    patience = 5,
    early_stopping = False,
    scheduler_factor = 0.5,
    scheduler_patience = 0,
    save_best_model = False,
    device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu'),
    weight_decay = 0,
```

)

```
# Initialize a new project
```

```
wandb.init(name = "Best_Model", project = 'dl22_HW6' , config = hyperparameters2)
```

Finishing last run (ID:1f1ci18c) before initializing another...

Waiting for W&B process to finish... **(success).**

0.001 MB of 0.009 MB uploaded (0.000 MB deduped)

Synced **cifar10_exp1**: https://wandb.ai/pranavshekhar2/dl22_HW6/runs/1f1ci18c

Synced 4 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: ./wandb/run-20221112_031825-1f1ci18c/logs

Successfully finished last run (ID:1f1ci18c). Initializing new run:

Tracking run with wandb version 0.13.5

Run data is saved locally in /content/wandb/run-20221112_031851-3nrs566m

Syncing run **Best_Model** to [Weights & Biases](#) ([docs](#))

Display W&B run

▼ Data Loaders, Loss Function, Optimizer

```
wandb.config = hyperparameters2
```

```
wandb.config
```

```
namespace(architecture='Cifar10_CustomNN_Experiment', batch_size=256,
dataset='Cifar10', device=device(type='cuda', index=0), early_stopping=False,
epochs=10,
file_model=PosixPath('/content/drive/MyDrive/Deep_Learning_UTD/Model/cifar10_custom_n
grad_clipping=True, learning_rate=0.01, log_batch=True, log_interval=100,
max_norm=1, momentum=0.9, output_dim=10, patience=5, save_best_model=False,
scheduler_factor=0.5, scheduler_patience=0, weight_decay=0)
```

```
# Fix seed value
```

```
SEED = 2345
```

```
random.seed(SEED)
```

```
np.random.seed(SEED)
```

```
torch.manual_seed(SEED)
```

```
torch.cuda.manual_seed(SEED)
```

```
torch.backends.cudnn.deterministic = True
```

```
# Data Loader
```

```
train_loader = torch.utils.data.DataLoader(trainset, batch_size=wandb.config.batch_size, s
```

```
valid_loader = torch.utils.data.DataLoader(validset, batch_size=wandb.config.batch_size, s
```

```
test_loader = torch.utils.data.DataLoader(testset, batch_size=wandb.config.batch_size, s
```

```
# model
```

```
model = Cifar10CNN_Experiment()
```



```
# Initialize weights from normal distribution with mean 0 and standard deviation 0.01
def init_weights(m):
    if type(m) == nn.Linear:
        torch.nn.init.kaiming_normal_(m.weight)
        torch.nn.init.zeros_(m.bias)

    if type(m) == nn.Conv2d:
        torch.nn.init.kaiming_normal_(m.weight)
        torch.nn.init.zeros_(m.bias)

model.to(wandb.config.device)
# model.apply(init_weights)

# loss_function
loss_function = nn.CrossEntropyLoss()

optimizer = torch.optim.SGD(model.parameters(),
                             lr = wandb.config.learning_rate,
                             momentum = wandb.config.momentum,
                             weight_decay = wandb.config.weight_decay)

scheduler = ReduceLROnPlateau(optimizer, mode='min', factor= wandb.config.scheduler_factor
                              patience=wandb.config.scheduler_patience, verbose=True)
```

▼ Training Model

```
wandb.watch(model, log = 'all', log_freq=25, log_graph=True)

wandb: logging graph, to disable use `wandb.watch(log_graph=False)`
[<wandb.wandb_torch.TorchGraph at 0x7f8fc1977d90>]

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

# See live graphs in the notebook.
%%wandb
batch_ct_train, batch_ct_valid = 0, 0
train_loss_history, train_acc_history, valid_loss_history, valid_acc_history = train_loop(
```

```
↳ Epoch : 1 / 10
Time to complete 1 is 0:00:16.345803
Learning rate: 0.01
Train Loss: 1.7244 | Train Accuracy: 39.0825%
Valid Loss: 1.4698 | Valid Accuracy: 48.9900%

Epoch : 2 / 10
Time to complete 2 is 0:00:16.486050
Learning rate: 0.01
Train Loss: 1.2539 | Train Accuracy: 55.7850%
Valid Loss: 1.1999 | Valid Accuracy: 57.5600%

Epoch : 3 / 10
Time to complete 3 is 0:00:19.379684
Learning rate: 0.01
Train Loss: 1.0662 | Train Accuracy: 62.5175%
Valid Loss: 1.0884 | Valid Accuracy: 61.8200%

Epoch : 4 / 10
Time to complete 4 is 0:00:16.196145
Learning rate: 0.01
Train Loss: 0.9564 | Train Accuracy: 66.7375%
Valid Loss: 1.0169 | Valid Accuracy: 64.2400%

Epoch : 5 / 10
Time to complete 5 is 0:00:17.563714
Learning rate: 0.01
Train Loss: 0.8727 | Train Accuracy: 69.6100%
Valid Loss: 1.0085 | Valid Accuracy: 64.4600%

Epoch : 6 / 10
Time to complete 6 is 0:00:16.294586
Learning rate: 0.01
Train Loss: 0.8015 | Train Accuracy: 72.1725%
Valid Loss: 0.9290 | Valid Accuracy: 67.3800%

Epoch : 7 / 10
Time to complete 7 is 0:00:16.183906
Learning rate: 0.01
Train Loss: 0.7437 | Train Accuracy: 74.2950%
Valid Loss: 0.9227 | Valid Accuracy: 67.8300%

Epoch 00008: reducing learning rate of group 0 to 5.0000e-03.
Epoch : 8 / 10
Time to complete 8 is 0:00:17.431813
Learning rate: 0.005
Train Loss: 0.6824 | Train Accuracy: 76.3450%
Valid Loss: 0.9295 | Valid Accuracy: 67.6100%

Epoch : 9 / 10
Time to complete 9 is 0:00:15.941199
Learning rate: 0.005
Train Loss: 0.5902 | Train Accuracy: 80.3350%
Valid Loss: 0.8890 | Valid Accuracy: 69.4000%

Epoch : 10 / 10
Time to complete 10 is 0:00:17.412580
Learning rate: 0.005
```

▼ Wandb Link

WANDB Link - https://wandb.ai/pranavshekhar2/dl22_HW6?workspace=user-pranavshekhar2

Colab paid products - [Cancel contracts here](#)

